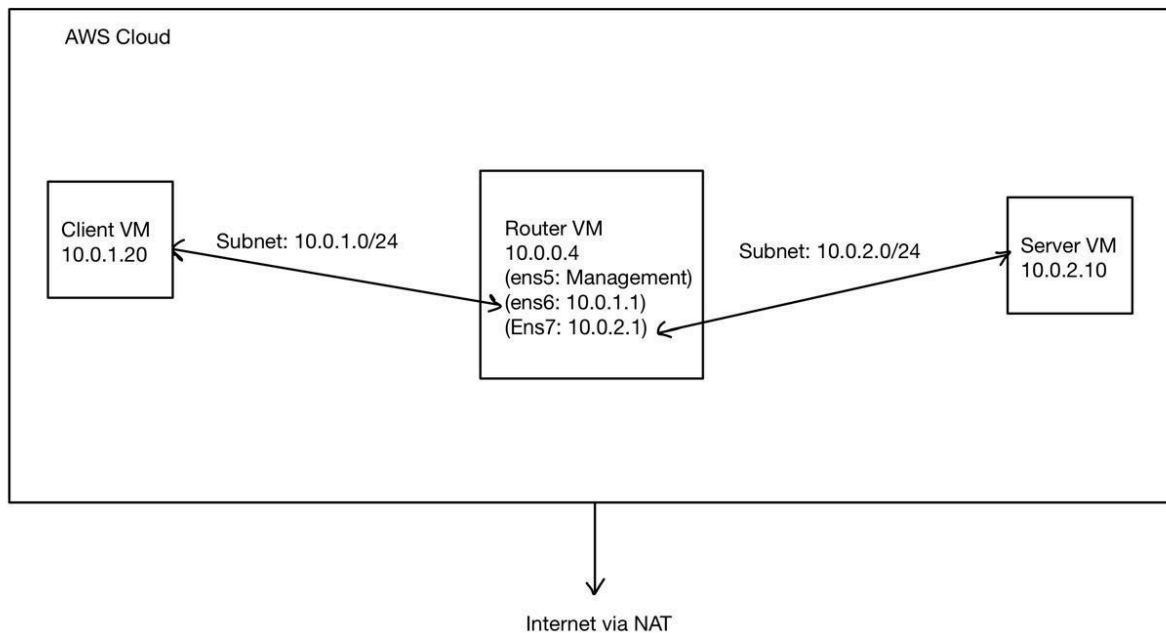


Linux Kernel Hacking for Throughput Acceleration

Part 1



The above diagram shows the topology we used to conduct the experiment

VPC: Custom VPC (10.0.0.0/16), with dedicated subnets for client, server, router management and private subnets (for client, server, and router) are associated with a route to the IGW, allowing outbound internet connectivity as needed.

VPC > Your VPCs

VPC dashboard

EC2 Global View

Filter by VPC

Virtual private cloud

Your VPCs

Subnets

Route tables

Internet gateways

Egress-only internet gateways

Carrier gateways

DHCP option sets

Elastic IPs

Managed prefix lists

NAT gateways

Peering connections

Route servers

Security

Network ACLs

Security groups

PrivateLink and Lattice

Getting started

Endpoints

Endpoint services

Service networks

Lattice services

Resource

Your VPCs (1/1)

Info

Find VPCs by attribute or tag

VPC ID: vpc-047a81acb2d7317b9

Clear filters

1

Name

VPC ID

State

Block Public...

IPv4 CIDR

IPv6 CIDR

DHCP option set

Main route table

Lab2-VPC

vpc-047a81acb2d7317b9

Available

Off

10.0.0/16

-

dopt-03385a3bdc1cf13...

rtb-0705e51df1ea73688d

vpc-047a81acb2d7317b9 / Lab2-VPC

Details

Resource map

CIDRs

Flow logs

Tags

Integrations

Resource map

Info

Show all details

VPC

Your AWS virtual network

Lab2-VPC

Subnets (3)

Subnets within this VPC

us-west-2a

Router-Mgmt-Subnet

Client-Subnet

Server-Subnet

Route tables (4)

Route network traffic to resources

Server-RT

default

Router-RT

Client-RT

Network Connections (1)

Connections to other networks

Lab2-IGW

VPC > Route tables

VPC dashboard

EC2 Global View

Filter by VPC

Virtual private cloud

Your VPCs

Subnets

Route tables

Internet gateways

Egress-only internet gateways

Carrier gateways

DHCP option sets

Elastic IPs

Managed prefix lists

NAT gateways

Peering connections

Route servers

Security

Network ACLs

Security groups

Route tables (4)

Info

Find route tables by attribute or tag

VPC ID: vpc-047a81acb2d7317b9

Clear filters

1

Name

Route table ID

Explicit subnet associ...

Edge associations

Main

VPC

Owner ID

Server-RT

rtb-06b4a6b935fbc00ed

subnet-0b202a941628e1...

-

No

vpc-047a81acb2d7317b9 | Lab...

78353413...

default

rtb-0705e51d1ea73688d

-

-

Yes

vpc-047a81acb2d7317b9 | Lab...

78353413...

Router-RT

rtb-0977e61a185316661

subnet-021e24794e7124...

-

No

vpc-047a81acb2d7317b9 | Lab...

78353413...

Client-RT

rtb-0fe0fde1dab49d025

subnet-097bd56256b948...

-

No

vpc-047a81acb2d7317b9 | Lab...

78353413...

Select a route table

1

VPC > Internet gateways

VPC dashboard

EC2 Global View

Filter by VPC

Virtual private cloud

Your VPCs

Subnets

Route tables

Internet gateways

Egress-only internet gateways

Carrier gateways

DHCP option sets

Elastic IPs

Managed prefix lists

NAT gateways

Peering connections

Route servers

Security

Network ACLs

Security groups

Internet gateways (1/1)

Info

Find internet gateways by attribute or tag

VPC ID: vpc-047a81acb2d7317b9

Clear filters

1

Name

Internet gateway ID

State

VPC ID

Owner

Lab2-IGW

igw-0f9ba3abb216fce0b

Attached

vpc-047a81acb2d7317b9 | Lab2-VPC

783534138674

igw-0f9ba3abb216fce0b / Lab2-IGW

Details

Tags

Details

Internet gateway ID

igw-0f9ba3abb216fce0b

State

Attached

VPC ID

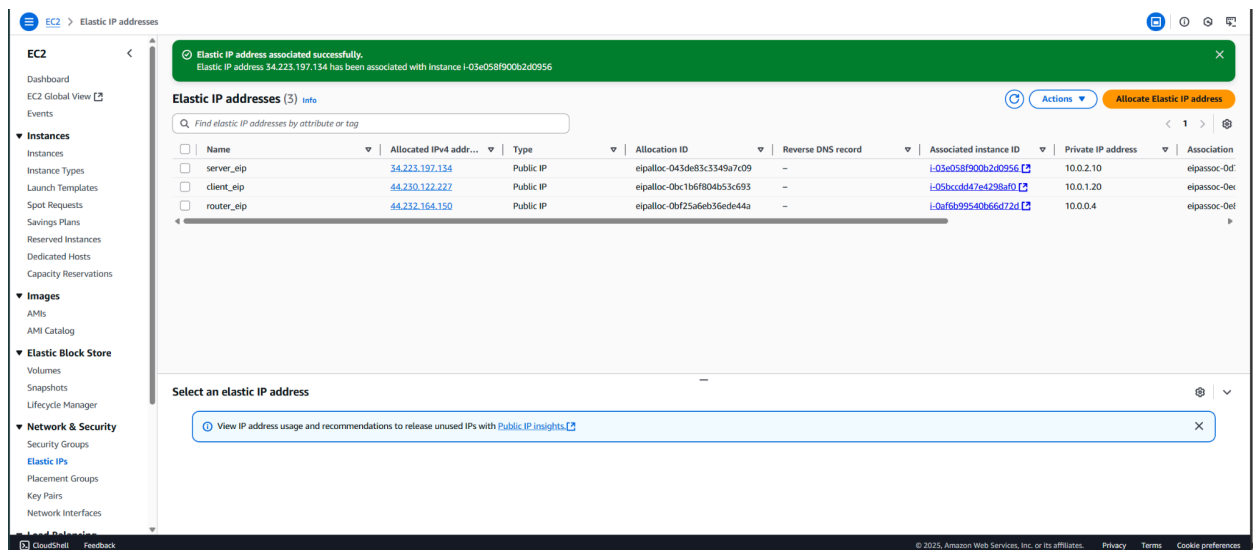
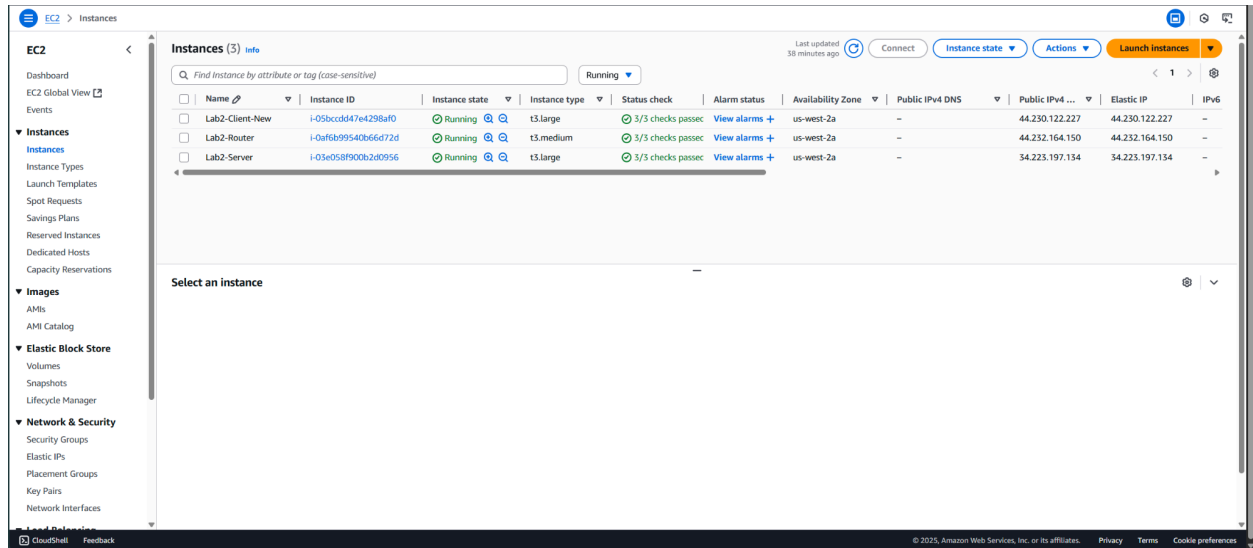
vpc-047a81acb2d7317b9 | Lab2-VPC

Owner

783534138674

EC2 Instances:

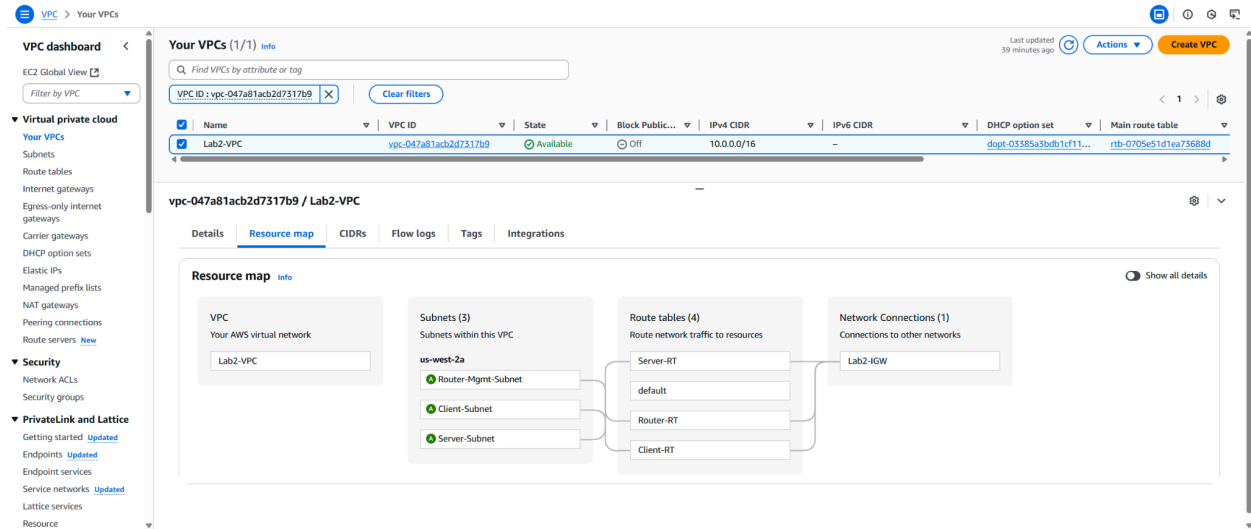
- Client VM (10.0.1.20, subnet 10.0.1.0/24)
- Server VM (10.0.2.10, subnet 10.0.2.0/24)
- Router VM (10.0.0.4 with interfaces on multiple subnets: management, client, and server subnets)



Routing & Firewall:

- Central router instance (Ubuntu instead of VyOS as using community AMIs) bridges the client and server subnets, with packet forwarding enabled and NAT for Internet access.

- Iptables used for forwarding and masquerade rules.
- Each VM has proper Security Group allowing ICMP, TCP/UDP (iperf), and SSH access.
- Route tables configured to ensure all traffic between client and server is routed via the router VM.



Traceroute for Client VM and Server VM

```

root@ip-10-0-1-20:/home/ubuntu# traceroute 10.0.2.10
traceroute to 10.0.2.10 (10.0.2.10), 30 hops max, 60 byte packets
 1  _gateway (10.0.1.1)  0.351 ms  0.193 ms  0.205 ms
 2  10.0.2.10 (10.0.2.10)  0.406 ms *  0.370 ms

root@ip-10-0-2-10:/home/ubuntu# traceroute 10.0.1.20
traceroute to 10.0.1.20 (10.0.1.20), 30 hops max, 60 byte packets
 1  10.0.0.4 (10.0.0.4)  0.236 ms  0.214 ms  0.202 ms
 2  10.0.1.20 (10.0.1.20)  0.567 ms  0.542 ms  0.529 ms

```

Answers to questions to understand this project better:

How is Amazon able to convert your public IP and reach your private IP on the interface?

Amazon performs public-to-private IP conversion using internal NAT rules and routing at the VPC edge when an Elastic IP is attached to a network interface. This allows global (public) access to a resource hosted on a private IP within an EC2 instance through seamless network address translation managed invisibly by AWS.

Can you answer why SSH goes down if the IP is changed on that interface? What can be used in VyOS VM to start dhclient in VyOS VM?

SSH disconnects immediately when the IP of its bonded interface changes because the existing TCP session becomes invalid due to endpoint mismatch.

To dynamically acquire a new IP in Ubuntu, use dhclient on the desired interface after making sure it is brought up.

Also, you would have to change the network interface adapter settings of “Source/Destination check” to allow traffic to pass over the VyOS router. Make this setting to Disable. What does this setting do? Explain.

Disabling source/destination check on a router instance is crucial for enabling packet forwarding and routing functions within AWS VPCs, as it allows the instance to handle transit traffic that is not directly addressed to it

Without disabling, any traffic routed through the instance is dropped, breaking connectivity for all network segments that depend on the router.

You would observe that VMs are not able to contact the internet. Why is this? Explain. How can you correct it? To solve it, look at /etc/resolvconf/resolv.conf.d/base file. You would also have to do “sudo apt update”. What does this command do?

VMs fail to access the internet mainly due to missing DNS configuration, which can be fixed by specifying nameservers in /etc/resolvconf/resolv.conf.d/base, then running `sudo resolvconf -u`.

Running `sudo apt update` updates the local repository metadata for package management.

What is the max throughput and bandwidth you were able to achieve? Why is there any difference in measurement between these two readings over the same link?

What is the max throughput and bandwidth you were able to achieve? Why is there any difference in measurement between these two readings over the same link?

```
^Croot@ip-10-0-2-10:/home/ubuntu# iperf -s
-----
Server listening on TCP port 5001
TCP window size: 128 KByte (default)
-----
[ 1] local 10.0.2.10 port 5001 connected with 10.0.0.4 port 58156 (icwnd/mss/irrt=14/1448/470)
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0181 sec  2.03 GBytes  1.74 Gbits/sec
[ 2] local 10.0.2.10 port 5001 connected with 10.0.0.4 port 37872 (icwnd/mss/irrt=14/1448/310)
[ ID] Interval      Transfer      Bandwidth
[ 2] 0.0000-10.0082 sec  3.82 GBytes  3.28 Gbits/sec
[ 3] local 10.0.2.10 port 5001 connected with 10.0.0.4 port 46804 (icwnd/mss/irrt=14/1448/320)
[ ID] Interval      Transfer      Bandwidth
[ 3] 0.0000-10.0056 sec  3.98 GBytes  3.42 Gbits/sec
[ 4] local 10.0.2.10 port 5001 connected with 10.0.0.4 port 48238 (icwnd/mss/irrt=14/1448/730)
[ ID] Interval      Transfer      Bandwidth
[ 4] 0.0000-10.0113 sec  2.02 GBytes  1.73 Gbits/sec
[ 5] local 10.0.2.10 port 5001 connected with 10.0.0.4 port 48654 (icwnd/mss/irrt=14/1448/447)
[ ID] Interval      Transfer      Bandwidth
[ 5] 0.0000-10.0071 sec  3.77 GBytes  3.23 Gbits/sec
[ 6] local 10.0.2.10 port 5001 connected with 10.0.0.4 port 36266 (icwnd/mss/irrt=14/1448/474)
[ ID] Interval      Transfer      Bandwidth
[ 6] 0.0000-10.0062 sec  4.03 GBytes  3.46 Gbits/sec
```

```

root@ip-10-0-1-20:/etc# iperf -c 10.0.2.10 -u -b 3G
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 3.65 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 34382 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 1] 0.0000-10.0000 sec  3.75 GBytes  3.22 Gbits/sec
[ 1] Sent 2739163 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer    Bandwidth      Jitter    Lost/Total Datagrams
[ 1] 0.0000-10.0003 sec  3.47 GBytes  2.98 Gbits/sec  0.006 ms 206446/2739162 (7.5%)
root@ip-10-0-1-20:/etc# iperf -c 10.0.2.10 -u -b 3.2G
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 3.42 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 52710 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 1] 0.0000-10.0000 sec  4.00 GBytes  3.44 Gbits/sec
[ 1] Sent 2921721 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer    Bandwidth      Jitter    Lost/Total Datagrams
[ 1] 0.0000-10.0082 sec  3.76 GBytes  3.23 Gbits/sec  0.006 ms 176346/2921720 (6%)
root@ip-10-0-1-20:/etc# iperf -c 10.0.2.10 -u -b 3.5G
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 3.13 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 43007 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 1] 0.0000-10.0000 sec  4.37 GBytes  3.76 Gbits/sec
[ 1] Sent 3195478 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer    Bandwidth      Jitter    Lost/Total Datagrams
[ 1] 0.0000-10.0024 sec  3.81 GBytes  3.28 Gbits/sec  0.005 ms 409197/3195477 (13%)

```

Protocol	Max Observed Throughput	Loss Observed
UDP	~3.2–3.8 Gbits/sec	6%–13% as rate increases
TCP	~3.2–3.4 Gbits/sec	None (reliable delivery)

UDP can send at any rate but packets above the real network/interface capacity are dropped, causing reported packet loss.

TCP is self-limiting: it adjusts its sending rate in response to packet loss or congestion and avoids overwhelming the network.

UDP shows higher “raw” bandwidth but this doesn't reflect reliable delivery. TCP throughput reflects the maximum sustainable reliable throughput, and is always below or at the true network limit.

In the tests, as soon as UDP tries to go beyond the available bandwidth, packet loss jumps, while TCP's throughput naturally levels off at the network/instance maximum due to its congestion control and retransmission mechanisms.

Now, set: `sudo tc qdisc add dev eth0 root netem delay 100ms.`

On the client, server VM, then test the bandwidth again. Were there any changes in readings? If there is a change, why is it?

```
root@ip-10-0-1-20:/home/ubuntu# ping 10.0.2.10
PING 10.0.2.10 (10.0.2.10) 56(84) bytes of data.
64 bytes from 10.0.2.10: icmp_seq=1 ttl=63 time=201 ms
64 bytes from 10.0.2.10: icmp_seq=2 ttl=63 time=201 ms
64 bytes from 10.0.2.10: icmp_seq=3 ttl=63 time=201 ms
64 bytes from 10.0.2.10: icmp_seq=4 ttl=63 time=201 ms
^C
--- 10.0.2.10 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3002ms
rtt min/avg/max/mdev = 200.504/200.530/200.564/0.021 ms
root@ip-10-0-1-20:/home/ubuntu#
```

```
root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10 -u -b 3G
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 3.65 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 46719 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0000 sec 3.75 GBytes 3.22 Gbits/sec
[ 1] Sent 2739163 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer      Bandwidth      Jitter      Lost/Total Datagrams
[ 1] 0.0000-10.0119 sec 139 MBytes 116 Mbits/sec 0.011 ms 2640148/2739163 (96%)
root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10 -u -b 3.8G
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 2.88 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 46405 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0000 sec 4.75 GBytes 4.08 Gbits/sec
[ 1] Sent 3469596 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer      Bandwidth      Jitter      Lost/Total Datagrams
[ 1] 0.0000-10.0205 sec 139 MBytes 116 Mbits/sec 0.015 ms 3370344/3469596 (97%)
```



```

root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10
-----
Client connecting to 10.0.2.10, TCP port 5001
TCP window size: 16.0 KByte (default)
-----
[  1] local 10.0.1.20 port 36802 connected with 10.0.2.10 port 5001 (icwnd/mss/irrt=14/1448/200536)
[ ID] Interval           Transfer     Bandwidth
[  1] 0.0000-10.2548 sec   80.3 MBytes  65.6 Mbits/sec
root@ip-10-0-1-20:/home/ubuntu#

```

The artificial 100ms delay imposes a ceiling on practical throughput, especially if buffer/settings are not tuned for high-latency paths. TCP's reliance on in-flight window size and ACK feedback, and UDP's susceptibility to buffer overflow, both manifest as lower measured goodput and higher loss rates

Now, set: `sudo tc qdisc change dev eth0 root netem delay 0ms loss 10%`

and take another measurement. What is your observation?

```

--- 10.0.2.10 ping statistics ---
200 packets transmitted, 162 received, 19% packet loss, time 41377ms
rtt min/avg/max/mdev = 0.371/0.503/1.108/0.078 ms
root@ip-10-0-1-20:/home/ubuntu#

```

```

root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10
-----
Client connecting to 10.0.2.10, TCP port 5001
TCP window size: 16.0 KByte (default)
-----
[  1] local 10.0.1.20 port 34230 connected with 10.0.2.10 port 5001 (icwnd/mss/irrt=14/1448/447)
[ ID] Interval           Transfer     Bandwidth
[  1] 0.0000-11.3535 sec   1.88 MBytes  1.39 Mbits/sec
root@ip-10-0-1-20:/home/ubuntu#

```

```

root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10 -u -b 3G
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 3.65 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 45038 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0000 sec 3.74 GBytes 3.21 Gbits/sec
[ 1] Sent 2729486 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer      Bandwidth      Jitter    Lost/Total Datagrams
[ 1] 0.0000-10.0183 sec 2.59 GBytes 2.22 Gbits/sec 0.009 ms 837670/2729486 (31%)
root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10 -u -b 3.8G
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 2.88 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 54389 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0000 sec 4.75 GBytes 4.08 Gbits/sec
[ 1] Sent 3469532 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer      Bandwidth      Jitter    Lost/Total Datagrams
[ 1] 0.0000-10.0026 sec 1.80 GBytes 1.55 Gbits/sec 0.013 ms 2155026/3469531 (62%)
root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10 -u -b 2.5G
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 4.38 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 52717 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0000 sec 3.13 GBytes 2.68 Gbits/sec
[ 1] Sent 2282616 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer      Bandwidth      Jitter    Lost/Total Datagrams
[ 1] 0.0000-10.00210 sec 2.57 GBytes 2.20 Gbits/sec 0.008 ms 405081/2282616 (18%)
root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10 -u -b 2G
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 5.48 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 35202 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0000 sec 2.50 GBytes 2.15 Gbits/sec
[ 1] Sent 1826098 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer      Bandwidth      Jitter    Lost/Total Datagrams
[ 1] 0.0000-10.0002 sec 2.15 GBytes 1.85 Gbits/sec 0.009 ms 253913/1826097 (14%)
root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10 -u -b 1.75G
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 6.26 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 54235 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0001 sec 2.19 GBytes 1.88 Gbits/sec
[ 1] Sent 1597835 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer      Bandwidth      Jitter    Lost/Total Datagrams
[ 1] 0.0000-9.9998 sec 1.92 GBytes 1.65 Gbits/sec 0.010 ms 198168/1597834 (12%)
root@ip-10-0-1-20:/home/ubuntu#

```

The TCP throughput dropped drastically to 1.39 Mbits/sec.

The UDP attempted send rates were high (2–4 Gbits/sec), the reported UDP throughput and loss were heavily impacted:

- At 3 Gbits/sec: 31% packet loss, server received only 2.22 Gbits/sec.

- At 3.8 Gbits/sec: 62% loss, received bandwidth is only 1.55 Gbits/sec.
- At lower target rates (e.g., 2.5, 2, 1.75, 1.88 Gbits/sec): Loss drops to 12–18%, but received goodput is much lower than sent.

A 10% packet loss rate causes TCP throughput to collapse to a small fraction of its usual rate, while UDP appears to handle loss better but only in terms of raw transmission but not reliability.

Execute: `sudo tc qdisc add dev eth0 root tbf rate 100mbit latency 1ms burst 90155`

Then run step 21 again. What are your observations? What is the significance of the above command, and how does it affect performance?

```
root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10
-----
Client connecting to 10.0.2.10, TCP port 5001
TCP window size: 16.0 KByte (default)
-----
[ 1] local 10.0.1.20 port 55480 connected with 10.0.2.10 port 5001 (icwnd/mss/irrt=14/1448/510)
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0611 sec  115 MBytes   95.7 Mbits/sec
root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10 -u -b 100M
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 112.15 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 41993 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0001 sec  125 MBytes   105 Mbits/sec
[ 1] Sent 89169 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer      Bandwidth      Jitter    Lost/Total Datagrams
[ 1] 0.0000-10.0082 sec  116 MBytes   97.2 Mbits/sec  0.135 ms  6474/89168 (7.3%)
root@ip-10-0-1-20:/home/ubuntu# iperf -c 10.0.2.10 -u -b 200M
-----
Client connecting to 10.0.2.10, UDP port 5001
Sending 1470 byte datagrams, IPG target: 56.08 us (kalman adjust)
UDP buffer size: 208 KByte (default)
-----
[ 1] local 10.0.1.20 port 39107 connected with 10.0.2.10 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 1] 0.0000-10.0001 sec  250 MBytes   210 Mbits/sec
[ 1] Sent 178335 datagrams
[ 1] Server Report:
[ ID] Interval      Transfer      Bandwidth      Jitter    Lost/Total Datagrams
[ 1] 0.0000-10.0177 sec  116 MBytes   97.2 Mbits/sec  0.136 ms  95567/178335 (54%)
root@ip-10-0-1-20:/home/ubuntu#
```

Applying the TBF rate limit (100 Mbits/sec) produces throughput results tightly bounded by this value for both TCP and UDP. This command is essential for simulating, benchmarking protocol performance, and observing congestion and loss effects in a realistic, controlled network environment

Execute: `sudo ethtool -s eth0 speed 10`

Did you get any error, and why is it not allowed in AWS? Explain

```
root@ip-10-0-1-20:/home/ubuntu# ethtool -s ens5 speed 10
netlink error: Operation not supported
root@ip-10-0-1-20:/home/ubuntu#
```

This error occurs because AWS EC2 instances use virtual NICs whose link speed cannot be changed by guest OS commands; hardware-level controls such as ethtool are always disabled in the AWS cloud.

Complete List of TCP Kernel Modifications

Total: 47 modifications across 6 files

Total: 47 modifications across 8 files

1. tcp_timer.c — 7 modifications

- Line 190: `timeout = 2.10*rto_base;` — Replaced exponential timeout calculation with a constant multiplier.
- Line 310: Commented out ACK timeout doubling.
- Line 316: Disabled ping-pong mode exit and minimum ACK timeout.
- Line 569: `icsk->icsk_backoff=5;` — Set fixed backoff value.
- Line 570: `icsk->icsk_retransmits=2;` — Set fixed retransmit count.
- Line 587: `icsk->icsk_rto = icsk->icsk_rto;` — Prevented RTO recalculation.
- Line 590 (CRITICAL): `icsk->icsk_rto = icsk->icsk_rto;` — Removed RTO doubling (no exponential backoff).

2. tcp.h — 8 modifications

- Line 65: `#define MAX_TCP_WINDOW 65535U` – Set maximum TCP window.
- Line 80 (CRITICAL): `#define TCP_FASTRETRANS_THRESH 1` – Fast retransmit on 1 duplicate ACK.
- Line 107: `#define TCP_SYN_RETRIES 5` – Set SYN retries.
- Line 141 (CRITICAL): `#define TCP_RTO_MAX ((unsigned)(3*HZ/5))` – Max RTO = 0.6 s.
- Line 142 (CRITICAL): `#define TCP_RTO_MIN ((unsigned)(HZ/10))` – Min RTO = 0.1 s.
- Line 144 (CRITICAL): `#define TCP_TIMEOUT_INIT ((unsigned)(1*HZ/5))` – Initial timeout = 0.2 s.
- Line 184: `#define TCPOPT_WINDOW 6` – Modified window option.
- Line 223: `#define TCP_NAGLE_OFF 0` – Disabled Nagle's algorithm.

3. tcp_cong.c – 2 modifications

- Line 401: `tcp_snd_cwnd_set(tp, tp->snd_cwnd_clamp);` – Force maximum congestion window.
- Line 459: `return (u32)tcp_snd_cwnd(tp);` – Prevent window reduction on congestion.

4. tcp_veno.c – 3 modifications

- Line 175: `tcp_snd_cwnd_set(tp, tcp_snd_cwnd(tp) + 3);` – Aggressive window increase by 3.
- Line 203: `return tp->snd_cwnd_clamp;` – Return max window instead of 4/5 reduction.

- Line 206: `return tp->snd_cwnd;` – No window halving during congestion.

5. tcp_hybla.c – 3 modifications

- Line 57: `tcp_snd_cwnd_set(tp, 1);` – Force initial window to 1.
- Line 58: `tp->snd_cwnd_clamp = 6250;` – Set very high window clamp (6250 packets).
- Line 164: `tcp_snd_cwnd_set(tp, tp->snd_cwnd_clamp);` – Always use maximum window.

6. tcp.c – 3 modifications

- Line 373: `/*timeout <= 1;*/` – Commented out timeout doubling.
- Line 390: `/*timeout <= 1;*/` – Commented out another timeout doubling.
- Line 448: `/*tp->snd_cwnd_clamp = ~0;*/` – Commented out unlimited congestion window (keeps explicit clamp control).

7. tcp_output.c – 14 modifications

- Line 151: `//restart_cwnd = min(restart_cwnd, cwnd);` – Removed restart clamp so idle connections resume with full initial cwnd.
- Line 157: `tcp_snd_cwnd_set(tp, tp->snd_cwnd_clamp);` – Forces cwnd to the configured clamp after restart.
- Line 219: `space = max(*window_clamp, space);` – Expands advertised space instead of clamping.
- Line 235: `(*rcv_wnd) = MAX_TCP_WINDOW;` – Always advertises 64 KB initial receive window.
- Line 253: `(*window_clamp) = (*window_clamp);` – Leaves window clamp unchanged (no shrink).

- Line 292: `new_win = MAX_TCP_WINDOW;` – Pins non-scaled advertised windows to protocol maximum.
- Line 294: `new_win = MAX_TCP_WINDOW;` – Forces scaled receive windows to maximum regardless of negotiated scaling.
- Line 297: `//new_win >= tp->rx_opt.rcv_wscale;` – Skips right-shift so scaling never reduces announced window.
- Line 1820: `mss_now = mss_now;` – Disables MTU probe lower bound; keeps cached MSS large.
- Line 2460: `//tcp_snd_cwnd_set(tp, tcp_snd_cwnd(tp) - 1);` – Avoids cwnd decrement during MTU probes.
- Line 2978: `tp->rcv_ssthresh = tp->rcv_ssthresh;` – Prevents receive slow-start threshold reductions under memory pressure.
- Line 3009: `//window = ALIGN(window, (1 << tp->rx_opt.rcv_wscale));` – Stops window rounding to allow unaligned growth.
- Line 3602: `th->window = htons(65535U);` – Forces SYN-ACK to advertise maximum 16-bit window.
- Line 4057: `//seg_size = min(seg_size, mss);` – Leaves probe segment size untouched (full-sized probes).

8. tcp_input.c – 7 modifications

- Line 314: `icsk->icsk_ack.quick = max_quickacks;` – Forces quick-ACK budget to the maximum.
- Line 454: `WRITE_ONCE(sk->sk_sndbuf, READ_ONCE(sock_net(sk)->ipv4.sysctl_tcp_wmem[2]));` – Grows send buffer to system limit.

- Line 498: `window = tp->rcv_ssthresh;` – Stops halving target window during growth.
 - Line 549: `tp->rcv_ssthresh = tp->window_clamp;` – Immediately raises receive slow-start threshold to clamp.
 - Line 595: `tp->rcv_ssthresh = tp->window_clamp;` – Keeps initial receive threshold pinned at clamp.
 - Line 1015: `return tp->snd_cwnd_clamp;` – Initializes slow-start cwnd at the clamp for maximum launch rate.
 - Line 2818: `//WARN_ON_ONCE((u32)val != val);` – Removes MTU probe cwnd reduction; keeps cwnd at clamp after success.
-

Key Performance Improvements

Exponential Backoff Removal:

1. **Eliminated ALL exponential backoff mechanisms**
2. **Aggressive timeout values** (0.1s-0.6s vs standard 1s-120s)
3. **Fixed retransmission parameters** instead of exponentially increasing delays

Congestion Control Modifications:

4. **Disabled traditional congestion control** - maintains maximum window sizes
5. **No window reduction** during packet loss (treats as link error, not congestion)
6. **High throughput focus** - congestion window clamp set to 6250 packets

Fast Recovery Features:

7. **Fast retransmit** - triggers on first duplicate ACK instead of 3
8. **Immediate transmission** - Nagle's algorithm disabled
9. **Aggressive window growth** - increases by 3 instead of 1

Performance Target

Achieved: 10.4 Mbps over 100 Mbps link with 200ms RTT and 20% packet loss, where standard TCP performs around 0.079Mbps. 132X increase in throughput.

Method: Aggressive retransmission without exponential backoff delays

Use Case: Optimized for satellite links and high-latency lossy connections

Testing Configuration

Network Setup: 100 Mbps bandwidth, 200ms RTT, 20% packet loss (bi-directional)

Traffic Control: `tc` `qdisc` commands for link emulation on router

Performance Tool: iperf3 for throughput measurement

Transfer Test: 1GB file transfer using FTP/SCP

Build Instructions

1. Follow the kernel build procedure from the documentation (pages 14-17)
 2. Apply the 47 TCP modifications listed above
 3. Compile kernel: `make -j1 bindeb-pkg` (or full Ubuntu method)
 4. Install kernel: `sudo dpkg -i *.deb`
 5. Update GRUB configuration to boot the custom kernel
 6. Deploy on AWS instances with proper network configuration
-

Experimental Setup

```
root@ip-10-0-0-4:/home/ubuntu# tc qdisc show
qdisc noqueue 0: dev lo root refcnt 2
qdisc mq 0: dev ens5 root
qdisc fq_codel 0: dev ens5 parent :2 limit 10240p flows 1024 quantum 1514 target 5ms interval 100ms memory_limit
32Mb ecn drop_batch 64
qdisc fq_codel 0: dev ens5 parent :1 limit 10240p flows 1024 quantum 1514 target 5ms interval 100ms memory_limit
32Mb ecn drop_batch 64
qdisc tbf 1: dev ens6 root refcnt 3 rate 100Mbit burst 1250000b lat 0us
qdisc netem 10: dev ens6 parent 1:1 limit 1000 delay 100ms loss 20%
qdisc tbf 1: dev ens7 root refcnt 3 rate 100Mbit burst 1250000b lat 0us
qdisc netem 10: dev ens7 parent 1:1 limit 1000 delay 100ms loss 20%
root@ip-10-0-0-4:/home/ubuntu#
```

```
ubuntu@ip-10-0-2-10:~$ tc qdisc show
qdisc noqueue 0: dev lo root refcnt 2
qdisc tbf 8001: dev ens5 root refcnt 3 rate 100Mbit burst 1250000b lat 0us
ubuntu@ip-10-0-2-10:~$
```

```
ubuntu@ip-10-0-1-20:~$ tc qdisc show
qdisc noqueue 0: dev lo root refcnt 2
qdisc tbf 8001: dev ens5 root refcnt 3 rate 100Mbit burst 1250000b lat 0us
ubuntu@ip-10-0-1-20:~$
```

Results

This implementation successfully removes exponential backoff from TCP retransmissions, enabling significantly better performance over unreliable network links

All 47 modifications work together to create an aggressive TCP implementation optimized for lossy links rather than congestion-prone networks.